

DATA WAREHOUSING AND DATA MINING PROJECT (LOW LEVEL IMPLEMENTATION THROUGH IPYNB NOTEBOOK)

CONTRIBUTORS

SACHIN SINGH - 2023BCS0064

ANSH RASTOGI - 2023BCS0151

ABHINAV VS - 2023BCS0181

TANISH BABU - 2023BCS0067

COVID-19 Data Warehousing and Data Mining Analysis

Objective

The objective of this project is to design a data warehousing pipeline and apply data mining techniques to analyze global COVID-19 trends. The project integrates ETL processes, OLAP operations, clustering analysis, outlier detection, association rule mining, and business intelligence visualizations to extract meaningful knowledge from pandemic data.

Dataset

Source: Our World in Data (OWID) COVID-19 Dataset

The dataset contains country-level pandemic indicators including cases, deaths, vaccinations, demographic and economic attributes.

Importing all the necessary modules for implementation

1. Importing Required Libraries

Libraries for:

- data manipulation (Pandas, NumPy)
- visualization (Matplotlib)
- clustering and mining (Scikit-learn)
- association rule mining

```
In [ ]: import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from scipy.cluster.hierarchy import linkage, dendrogram
from scipy.stats import zscore

from mlxtend.frequent_patterns import apriori, association_rules
```

2. Data Loading and Understanding

The dataset is loaded into a dataframe to understand structure, attributes, and completeness before preprocessing.

```
In [ ]: import pandas as pd

df = pd.read_csv("owid-covid-data.csv")
df.head()
```

```
Out[ ]:   iso_code continent    location      date  total_cases  new_cases  new_case
```

	iso_code	continent	location	date	total_cases	new_cases	new_case
0	AFG	Asia	Afghanistan	2020-01-05	0.0	0.0	0.0
1	AFG	Asia	Afghanistan	2020-01-06	0.0	0.0	0.0
2	AFG	Asia	Afghanistan	2020-01-07	0.0	0.0	0.0
3	AFG	Asia	Afghanistan	2020-01-08	0.0	0.0	0.0
4	AFG	Asia	Afghanistan	2020-01-09	0.0	0.0	0.0

5 rows × 67 columns

Dataset Overview

We inspect attribute types and missing values to prepare for cleaning and transformation.

```
In [ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 429435 entries, 0 to 429434
Data columns (total 67 columns):
 #   Column           Non-Null Count Dtype
 --- 
 0   iso_code         429435 non-null object
 1   continent        402910 non-null object
 2   location          429435 non-null object
 3   date              429435 non-null object
 4   total_cases       411804 non-null float64
 5   new_cases         410159 non-null float64
 6   new_cases_smoothed 408929 non-null float64
 7   total_deaths      411804 non-null float64
 8   new_deaths        410608 non-null float64
 9   new_deaths_smoothed 409378 non-null float64
 10  total_cases_per_million 411804 non-null float64
 11  new_cases_per_million 410159 non-null float64
 12  new_cases_smoothed_per_million 408929 non-null float64
 13  total_deaths_per_million 411804 non-null float64
 14  new_deaths_per_million 410608 non-null float64
 15  new_deaths_smoothed_per_million 409378 non-null float64
 16  reproduction_rate 184817 non-null float64
 17  icu_patients      39116 non-null float64
 18  icu_patients_per_million 39116 non-null float64
 19  hosp_patients     40656 non-null float64
 20  hosp_patients_per_million 40656 non-null float64
 21  weekly_icu_admissions 10993 non-null float64
 22  weekly_icu_admissions_per_million 10993 non-null float64
 23  weekly_hosp_admissions 24497 non-null float64
 24  weekly_hosp_admissions_per_million 24497 non-null float64
 25  total_tests        79387 non-null float64
 26  new_tests          75403 non-null float64
 27  total_tests_per_thousand 79387 non-null float64
 28  new_tests_per_thousand 75403 non-null float64
 29  new_tests_smoothed 103965 non-null float64
 30  new_tests_smoothed_per_thousand 103965 non-null float64
 31  positive_rate      95927 non-null float64
 32  tests_per_case     94348 non-null float64
 33  tests_units         106788 non-null object
 34  total_vaccinations 85417 non-null float64
 35  people_vaccinated 81132 non-null float64
 36  people_fully_vaccinated 78061 non-null float64
 37  total_boosters      53600 non-null float64
 38  new_vaccinations    70971 non-null float64
 39  new_vaccinations_smoothed 195029 non-null float64
 40  total_vaccinations_per_hundred 85417 non-null float64
 41  people_vaccinated_per_hundred 81132 non-null float64
 42  people_fully_vaccinated_per_hundred 78061 non-null float64
 43  total_boosters_per_hundred 53600 non-null float64
 44  new_vaccinations_smoothed_per_million 195029 non-null float64
 45  new_people_vaccinated_smoothed 192177 non-null float64
 46  new_people_vaccinated_smoothed_per_hundred 192177 non-null float64
 47  stringency_index    196190 non-null float64
 48  population_density 360492 non-null float64

```

```

49 median_age           334663 non-null float64
50 aged_65_older        323270 non-null float64
51 aged_70_older        331315 non-null float64
52 gdp_per_capita       328292 non-null float64
53 extreme_poverty      211996 non-null float64
54 cardiovasc_death_rate 328865 non-null float64
55 diabetes_prevalence 345911 non-null float64
56 female_smokers       247165 non-null float64
57 male_smokers         243817 non-null float64
58 handwashing_facilities 161741 non-null float64
59 hospital_beds_per_thousand 290689 non-null float64
60 life_expectancy       390299 non-null float64
61 human_development_index 319127 non-null float64
62 population            429435 non-null int64
63 excess_mortality_cumulative_absolute 13411 non-null float64
64 excess_mortality_cumulative           13411 non-null float64
65 excess_mortality             13411 non-null float64
66 excess_mortality_cumulative_per_million 13411 non-null float64
dtypes: float64(61), int64(1), object(5)
memory usage: 219.5+ MB

```

3. Attribute Selection

Relevant attributes were selected based on pandemic impact and response indicators.

This reduces dimensionality and improves mining efficiency.

```
In [ ]: cols = [
    'location',
    'continent',
    'date',
    'total_cases',
    'total_deaths',
    'people_vaccinated',
    'population',
    'gdp_per_capita'
]

df = df[cols]
```

4. Data Cleaning

Cleaning removes inconsistencies:

- rows without continent information removed
- missing numerical values replaced
- date converted to datetime format

```
In [ ]: df = df.dropna(subset=['continent'])

In [ ]: df[['total_cases','total_deaths','people_vaccinated']] = \
df[['total_cases','total_deaths','people_vaccinated']].fillna(0)

In [ ]: df['date'] = pd.to_datetime(df['date'])
```

5. Data Transformation (Feature Engineering)

New analytical indicators are created:

- Death Rate → severity indicator
- Vaccination Rate → healthcare response effectiveness

```
In [ ]: df['death_rate'] = (
    df['total_deaths'] / df['total_cases']
).replace([np.inf, -np.inf],0).fillna(0)

df['vaccination_rate'] = (
    df['people_vaccinated'] / df['population']
).replace([np.inf, -np.inf],0).fillna(0)
```

```
In [ ]: df[['death_rate','vaccination_rate']].head()
```

```
Out[ ]:   death_rate  vaccination_rate
0          0.0           0.0
1          0.0           0.0
2          0.0           0.0
3          0.0           0.0
4          0.0           0.0
```

6. Data Discretization and Concept Hierarchy

Population values are categorized into levels to support multidimensional OLAP analysis.

```
In [ ]: def pop_category(x):
    if x < 1e7:
        return "Low"
    elif x < 1e8:
        return "Medium"
    else:
        return "High"
```

```
df['pop_category'] = df['population'].apply(pop_category)
```

```
In [ ]: reduced_df = df.copy()
```

7. Data Warehouse Construction

A Star Schema is implemented consisting of:

- Country Dimension
- Date Dimension
- Fact Table containing analytical measures.

```
In [ ]: dim_country = reduced_df[['location','continent','population','gdp_per_capita']].drop_duplicates()  
dim_country['country_id'] = range(len(dim_country))
```

```
In [ ]: dim_date = reduced_df[['date']].drop_duplicates()  
  
dim_date['year'] = dim_date['date'].dt.year  
dim_date['month'] = dim_date['date'].dt.month  
dim_date['quarter'] = dim_date['date'].dt.quarter
```

```
In [ ]: fact_table = reduced_df.merge(  
        dim_country[['location','country_id']],  
        on='location'  
)
```

```
In [ ]: fact_table.columns
```

```
Out[ ]: Index(['location', 'continent', 'date', 'total_cases', 'total_deaths',  
              'people_vaccinated', 'population', 'gdp_per_capita', 'death_rate',  
              'vaccination_rate', 'pop_category', 'country_id'],  
              dtype='object')
```

8. OLAP Operations

Multidimensional analysis performed using:

- Roll-Up
- Drill-Down
- Slice
- Dice

```
In [ ]: monthly_cases = fact_table.groupby(
```

```
['location', fact_table['date'].dt.to_period('M')]  
)[['total_cases']].sum().reset_index()
```

```
In [ ]: continent_country = fact_table.groupby(  
['continent', 'location'])  
['total_cases'].sum()
```

```
In [ ]: dice_data = fact_table[  
(fact_table['continent'].isin(['Asia', 'Europe'])) &  
(fact_table['date'].dt.year == 2021)  
]
```

9. Data Cube Representation

A cube is created using Continent and Year dimensions with Total Cases as measure.

```
In [ ]: cube = fact_table.pivot_table(  
values='total_cases',  
index='continent',  
columns=fact_table['date'].dt.year,  
aggfunc='sum'  
)  
  
cube
```

```
Out[ ]:      date      2020      2021      2022      2023      2024  
continent  
  Africa  2.850575e+08  2.222352e+09  4.480156e+09  4.777568e+09  2.851753e+09  
  Asia    2.032098e+09  1.945136e+10  5.678968e+10  1.084923e+11  6.540192e+10  
  Europe   1.625730e+09  1.815965e+10  7.157639e+10  9.061265e+10  5.478227e+10  
  North America 1.952655e+09  1.545812e+10  3.751653e+10  4.513684e+10  2.700953e+10  
  Oceania   6.973459e+06  5.865609e+07  3.132067e+09  5.173457e+09  3.226900e+09  
  South America 1.456051e+09  1.077255e+10  2.139144e+10  2.494248e+10  1.492205e+10
```

10. Global Pandemic Trend Analysis


This visualization shows how COVID cases evolved over time globally.

Insight

Multiple peaks indicate pandemic waves and policy response cycles.

```
In [ ]: global_cases = fact_table.groupby('date')['total_cases'].sum()

plt.figure()
plt.plot(global_cases)
plt.title("Global COVID Cases Trend")
plt.show()
```



11. Vaccination Effectiveness Analysis

Examines relationship between vaccination coverage and mortality.

Observation

Higher vaccination rates generally correspond to lower death rates.

```
In [96]: plt.figure(figsize=(8,5))


plt.scatter(
fact_table['vaccination_rate'],
```

```

fact_table['death_rate']
)

plt.xlabel("Vaccination Rate")
plt.ylabel("Death Rate")
plt.show()

```



12. K-Means Clustering (Partitioning Method)

Countries are grouped based on pandemic similarity using:

- total cases
- total deaths
- vaccination rate

```
In [ ]: cluster_data = fact_table.groupby('location')[[
    'total_cases',
    'total_deaths',
    'vaccination_rate'
]].mean()
```

```
In [ ]: scaler = StandardScaler()
X = scaler.fit_transform(cluster_data)
```

```
In [ ]: kmeans = KMeans(n_clusters=3, random_state=0)
cluster_data['cluster'] = kmeans.fit_predict(X)
```

```
In [ ]: cluster_sample = cluster_data.sort_values(  
    'total_cases',  
    ascending=False  
) .head(30)
```


```
In [ ]: from scipy.cluster.hierarchy import linkage, dendrogram  
  
Z = linkage(  
    cluster_sample[['total_cases', 'total_deaths', 'vaccination_rate']],  
    method='ward'  
)
```

Cluster Visualization

Each point represents a country. Colors indicate cluster membership.

Clusters reveal groups of countries with similar pandemic impact levels.

```
In [97]: plt.figure(figsize=(8,5))  
  
plt.scatter(  
    cluster_data['total_cases'],  
    cluster_data['vaccination_rate'],  
    c=cluster_data['cluster']  
)  
  
plt.xlabel("Total Cases")  
plt.ylabel("Vaccination Rate")  
plt.show()
```



13. Hierarchical Clustering


A dendrogram visualizes similarity relationships among countries without predefined cluster numbers.

```
In [ ]: plt.figure(figsize=(12,6))

dendrogram(
Z,
labels=cluster_sample.index,
leaf_rotation=90
)

plt.title("Hierarchical Clustering of Top 30 Countries")
plt.xlabel("Countries")
plt.ylabel("Distance")

plt.show()
```



14. Outlier Detection

Z-score analysis identifies countries with abnormal pandemic patterns.

```
In [ ]: cluster_data['zscore_cases'] = zscore(cluster_data['total_cases'])

outliers = cluster_data[
abs(cluster_data['zscore_cases']) > 3
]

outliers
```

	total_cases	total_deaths	vaccination_rate	cluster	zscore_cases	
location						
Brazil	2.405882e+07	505097.671446		0.257536	1	3.759158
China	3.571664e+07	47358.838710		0.024245	1	5.724301
France	2.107340e+07	117994.496416		0.365552	1	3.255909
Germany	1.994243e+07	113734.657706		0.310501	1	3.065263
India	3.096283e+07	378511.336504		0.437517	1	4.922958
United States	6.327030e+07	777909.996416		0.331806	1	10.368985

```
In [ ]: plt.figure(figsize=(10,5))


plt.scatter(
range(len(cluster_data)),
cluster_data['total_cases'],
label="Normal Data"
)

outliers = cluster_data[abs(cluster_data['zscore_cases'])] > 3

plt.scatter(
outliers.index.map(lambda x: cluster_data.index.get_loc(x)),
outliers['total_cases'],
label="Outliers"
)

plt.title("Outlier Detection using Z-Score")
plt.xlabel("Countries (index)")
plt.ylabel("Average Total Cases")

plt.legend()
plt.show()
```



Continent-Level Pandemic Impact Analysis (OLAP Roll-Up Visualization)

This visualization presents the total number of COVID-19 cases aggregated at the continent level.

The aggregation represents an **OLAP Roll-Up operation**, where detailed country-

level data is summarized into a higher-level geographical dimension.

Purpose

- To compare pandemic impact across continents.
- To identify regions contributing most to global case counts.
- To demonstrate multidimensional aggregation using data warehousing concepts.

Analytical Observations


By aggregating cases by continent, large-scale regional trends become visible that are not easily observable at the country level. Differences in total cases may reflect variations in population size, mobility patterns, healthcare infrastructure, and policy responses.

Insight

This visualization supports Business Intelligence by enabling decision-makers to quickly assess regional pandemic severity and prioritize resource allocation at a macro level.

```
In [ ]: continent_cases = fact_table.groupby('continent')['total_cases'].sum()

plt.figure()
continent_cases.plot(kind='bar')
plt.title("Total Cases by Continent")
plt.ylabel("Cases")
plt.show()
```



Top 10 Countries by Total COVID-19 Cases (Business Intelligence Visualization)

This visualization identifies the ten countries with the highest recorded total COVID-19 cases.

The data is aggregated at the country level by selecting the maximum reported case count for each location, representing the peak pandemic impact experienced by each country.

Purpose

- To highlight the most affected countries globally.
- To support comparative analysis between nations.
- To provide a clear Business Intelligence view of pandemic severity.

Analytical Observations


The horizontal bar chart enables easy comparison of countries with extreme case counts. Countries appearing in this ranking typically have large populations, high urban density, or experienced multiple infection waves.

Insight

This visualization helps decision-makers quickly identify regions requiring significant healthcare resources and policy intervention. It also demonstrates how aggregation transforms raw data into actionable insights for strategic planning.

```
In [ ]: top10 = fact_table.groupby('location')['total_cases'].max().nlargest(10)

plt.figure()
top10.plot(kind='barh')
plt.title("Top 10 Countries by Total Cases")
plt.xlabel("Cases")
plt.show()
```



Multidimensional Data Cube Visualization (Continent vs Year)

This visualization represents a **Data Cube** constructed using multidimensional aggregation.

The cube summarizes total COVID-19 cases across two analytical dimensions:

- **Geographical Dimension:** Continent
- **Temporal Dimension:** Year
- **Measure:** Total COVID-19 Cases

The pivot table performs aggregation similar to OLAP cube construction, and the heatmap provides a visual representation of intensity across dimensions.

Purpose

- To demonstrate multidimensional data representation.
- To analyze how pandemic impact varies across regions and time.
- To support OLAP-based analytical exploration.

Analytical Observations

Color intensity reflects the magnitude of total cases. Darker regions indicate periods and continents experiencing higher infection levels. The visualization makes temporal trends and regional disparities immediately visible without inspecting numerical tables.

Insight


The data cube enables decision-makers to compare pandemic progression across continents over multiple years simultaneously. This illustrates how data warehousing techniques transform large datasets into structured analytical views suitable for Business Intelligence and strategic analysis.

```
In [ ]: cube = fact_table.pivot_table(  
values='total_cases',  
index='continent',  
columns=fact_table['date'].dt.year,  
aggfunc='sum'  
)  
  
plt.figure()  
plt.imshow(cube)  
plt.colorbar()
```

```

plt.xticks(range(len(cube.columns)), cube.columns)
plt.yticks(range(len(cube.index)), cube.index)
plt.title("Data Cube Heatmap (Continent vs Year)")
plt.show()

```



Vaccination Rate Distribution Analysis

This histogram visualizes the distribution of vaccination rates across countries in the dataset.

The purpose of this analysis is to understand how vaccination coverage is spread globally and to identify common patterns, concentration ranges, and variability among nations.

Purpose

- To examine the statistical distribution of vaccination rates.
- To understand data characteristics before applying mining techniques.
- To identify whether vaccination adoption is uniform or uneven across countries.


Analytical Observations

The histogram shows how frequently different vaccination rate ranges occur. Peaks in specific intervals indicate where most countries fall in terms of vaccination coverage, while sparse regions represent countries with unusually low or high vaccination levels.

Insight

The distribution reveals global inequality in vaccination rollout, where many countries cluster around moderate vaccination levels while fewer achieve extremely high coverage. Understanding this distribution helps interpret clustering results and supports further analysis of vaccination effectiveness.

```
In [ ]: plt.figure()  
plt.hist(fact_table['vaccination_rate'], bins=30)  
plt.title("Vaccination Rate Distribution")  
plt.show()
```



16. Knowledge Extraction and Conclusion

Key findings:

- Vaccination significantly reduces mortality risk.
- Countries form natural clusters based on pandemic severity.
- Outlier countries show abnormal spread characteristics.
- OLAP and data cubes enable multidimensional analysis.

The project demonstrates how data warehousing combined with data mining converts raw data into actionable intelligence.